

Propositions

- An interpretation is an assignment of values to all variables.
- A model is an interpretation that satisfies the constraints.
- Often we don't want to just find a model, but want to know what is true in all models.
- A proposition is statement that is true or false in each interpretation.

Why propositions?

- Specifying logical formulae is often more natural than filling in tables
- It is easier to check correctness and debug formulae than tables
- We can exploit the Boolean nature for efficient reasoning
- We need a language for asking queries (of what follows in all models) that may be more complicated than asking for the value of a variable
- It is easy to incrementally add formulae
- It can be extended to infinitely many variables with infinite domains (using logical quantification)

Human's view of semantics

Step 1 Begin with a task domain.

Step 2 Choose atoms in the computer to denote propositions. These atoms have meaning to the KB designer.

Step 3 Tell the system knowledge about the domain.

Step 4 Ask the system questions.

— the system can tell you whether the question is a logical consequence.

— You can interpret the answer with the meaning associated with the atoms.

In computer:

$light1_broken \leftarrow sw_up$
 $\wedge power \wedge unlit_light1.$
 $sw_up.$
 $power \leftarrow lit_light2.$
 $unlit_light1.$
 $lit_light2.$

In user's mind:

- $light1_broken$: light #1 is broken
- sw_up : switch is up
- $power$: there is power in the building
- $unlit_light1$: light #1 isn't lit
- lit_light2 : light #2 is lit

Conclusion: $light1_broken$

- The computer doesn't know the meaning of the symbols
- The user can interpret the symbol using their meaning

Simple language: propositional definite clauses

- An **atom** is a symbol starting with a lower case letter
- A **body** is an atom or is of the form $b_1 \wedge b_2$ where b_1 and b_2 are bodies.
- A **definite clause** is an atom or is a rule of the form $h \leftarrow b$ where h is an atom and b is a body.
- A **knowledge base** is a set of definite clauses

- An **interpretation** I assigns a truth value to each atom.
- A body $b_1 \wedge b_2$ is true in I if b_1 is true in I and b_2 is true in I .
- A rule $h \leftarrow b$ is false in I if b is true in I and h is false in I .
The rule is true otherwise.
- A knowledge base KB is true in I if and only if every clause in KB is true in I .

Models and Logical Consequence

- A **model** of a set of clauses is an interpretation in which all the clauses are *true*.
- If KB is a set of clauses and g is a conjunction of atoms, g is a **logical consequence** of KB , written $KB \models g$, if g is *true* in every model of KB .
- That is, $KB \models g$ if there is no interpretation in which KB is *true* and g is *false*.

Simple Example

$$KB = \begin{cases} p \leftarrow q. \\ q. \\ r \leftarrow s. \end{cases}$$

| | <i>p</i> | <i>q</i> | <i>r</i> | <i>s</i> | model? |
|-----------------------|--------------|--------------|--------------|--------------|--------|
| <i>l</i> ₁ | <i>true</i> | <i>true</i> | <i>true</i> | <i>true</i> | |
| <i>l</i> ₂ | <i>false</i> | <i>false</i> | <i>false</i> | <i>false</i> | |
| <i>l</i> ₃ | <i>true</i> | <i>true</i> | <i>false</i> | <i>false</i> | |
| <i>l</i> ₄ | <i>true</i> | <i>true</i> | <i>true</i> | <i>false</i> | |
| <i>l</i> ₅ | <i>true</i> | <i>true</i> | <i>false</i> | <i>true</i> | |

Simple Example

$$KB = \begin{cases} p \leftarrow q. \\ q. \\ r \leftarrow s. \end{cases}$$

| | <i>p</i> | <i>q</i> | <i>r</i> | <i>s</i> | model? |
|-----------------------|--------------|--------------|--------------|--------------|--------------------------|
| <i>l</i> ₁ | <i>true</i> | <i>true</i> | <i>true</i> | <i>true</i> | is a model of <i>KB</i> |
| <i>l</i> ₂ | <i>false</i> | <i>false</i> | <i>false</i> | <i>false</i> | not a model of <i>KB</i> |
| <i>l</i> ₃ | <i>true</i> | <i>true</i> | <i>false</i> | <i>false</i> | is a model of <i>KB</i> |
| <i>l</i> ₄ | <i>true</i> | <i>true</i> | <i>true</i> | <i>false</i> | is a model of <i>KB</i> |
| <i>l</i> ₅ | <i>true</i> | <i>true</i> | <i>false</i> | <i>true</i> | not a model of <i>KB</i> |

Simple Example

$$KB = \begin{cases} p \leftarrow q. \\ q. \\ r \leftarrow s. \end{cases}$$

| | <i>p</i> | <i>q</i> | <i>r</i> | <i>s</i> | model? |
|-----------------------|--------------|--------------|--------------|--------------|--------------------------|
| <i>l</i> ₁ | <i>true</i> | <i>true</i> | <i>true</i> | <i>true</i> | is a model of <i>KB</i> |
| <i>l</i> ₂ | <i>false</i> | <i>false</i> | <i>false</i> | <i>false</i> | not a model of <i>KB</i> |
| <i>l</i> ₃ | <i>true</i> | <i>true</i> | <i>false</i> | <i>false</i> | is a model of <i>KB</i> |
| <i>l</i> ₄ | <i>true</i> | <i>true</i> | <i>true</i> | <i>false</i> | is a model of <i>KB</i> |
| <i>l</i> ₅ | <i>true</i> | <i>true</i> | <i>false</i> | <i>true</i> | not a model of <i>KB</i> |

Which of *p*, *q*, *r*, *q* logically follow from *KB*?

Simple Example

$$KB = \begin{cases} p \leftarrow q. \\ q. \\ r \leftarrow s. \end{cases}$$

| | <i>p</i> | <i>q</i> | <i>r</i> | <i>s</i> | model? |
|-----------------------|--------------|--------------|--------------|--------------|--------------------------|
| <i>l</i> ₁ | <i>true</i> | <i>true</i> | <i>true</i> | <i>true</i> | is a model of <i>KB</i> |
| <i>l</i> ₂ | <i>false</i> | <i>false</i> | <i>false</i> | <i>false</i> | not a model of <i>KB</i> |
| <i>l</i> ₃ | <i>true</i> | <i>true</i> | <i>false</i> | <i>false</i> | is a model of <i>KB</i> |
| <i>l</i> ₄ | <i>true</i> | <i>true</i> | <i>true</i> | <i>false</i> | is a model of <i>KB</i> |
| <i>l</i> ₅ | <i>true</i> | <i>true</i> | <i>false</i> | <i>true</i> | not a model of <i>KB</i> |

Which of *p*, *q*, *r*, *s* logically follow from *KB*?

$KB \models p$, $KB \models q$, $KB \not\models r$, $KB \not\models s$


User's view of Semantics

1. Choose a task domain: **intended interpretation.**
2. Associate an atom with each proposition you want to represent.
3. Tell the system clauses that are true in the intended interpretation: **axiomatizing the domain.**
4. Ask questions about the intended interpretation.
5. If $KB \models g$, then g must be true in the intended interpretation.
6. Users can interpret the answer using their intended interpretation of the symbols.

Computer's view of semantics

- The computer doesn't have access to the intended interpretation.
- All it knows is the knowledge base.
- The computer can determine if a formula is a logical consequence of KB.
- If $KB \models g$ then g must be true in the intended interpretation.
- If $KB \not\models g$ then there is a model of KB in which g is false. This could be the intended interpretation.

Electrical Environment



Representing the Electrical Environment

light_l1.

light_l2.

down_s1.

up_s2.

up_s3.

ok_l1.

ok_l2.

ok_cb1.

ok_cb2.

live_outside.

lit_l1 \leftarrow *live_w0* \wedge *ok_l1*

live_w0 \leftarrow *live_w1* \wedge *up_s2.*

live_w0 \leftarrow *live_w2* \wedge *down_s2.*

live_w1 \leftarrow *live_w3* \wedge *up_s1.*

live_w2 \leftarrow *live_w3* \wedge *down_s1.*

lit_l2 \leftarrow *live_w4* \wedge *ok_l2.*

live_w4 \leftarrow *live_w3* \wedge *up_s3.*

live_p1 \leftarrow *live_w3.*

live_w3 \leftarrow *live_w5* \wedge *ok_cb1.*

live_p2 \leftarrow *live_w6.*

live_w6 \leftarrow *live_w5* \wedge *ok_cb2.*

live_w5 \leftarrow *live_outside.*

- A **proof** is a mechanically derivable demonstration that a formula logically follows from a knowledge base.
- Given a proof procedure, $KB \vdash g$ means g can be derived from knowledge base KB .
- Recall $KB \models g$ means g is true in all models of KB .
- A proof procedure is **sound** if $KB \vdash g$ implies $KB \models g$.
- A proof procedure is **complete** if $KB \models g$ implies $KB \vdash g$.

Bottom-up Ground Proof Procedure

One **rule of derivation**, a generalized form of *modus ponens*:

If " $h \leftarrow b_1 \wedge \dots \wedge b_m$ " is a clause in the knowledge base, and each b_i has been derived, then h can be derived.

This is **forward chaining** on this clause.

(This rule also covers the case when $m = 0$.)

Bottom-up proof procedure

$KB \vdash g$ if $g \in C$ at the end of this procedure:

$C := \{\}$;

repeat

select clause " $h \leftarrow b_1 \wedge \dots \wedge b_m$ " in KB such that

$b_i \in C$ for all i , and

$h \notin C$;

$C := C \cup \{h\}$

until no more clauses can be selected.

Example

$$a \leftarrow b \wedge c.$$

$$a \leftarrow e \wedge f.$$

$$b \leftarrow f \wedge k.$$

$$c \leftarrow e.$$

$$d \leftarrow k.$$

$$e.$$

$$f \leftarrow j \wedge e.$$

$$f \leftarrow c.$$

$$j \leftarrow c.$$

Soundness of bottom-up proof procedure

If $KB \vdash g$ then $KB \models g$.

- Suppose there is a g such that $KB \vdash g$ and $KB \not\models g$.
- Then there must be a first atom added to C that isn't true in every model of KB . Call it h . Suppose h isn't true in model I of KB .
- There must be a clause in KB of form

$$h \leftarrow b_1 \wedge \dots \wedge b_m$$

Each b_i is true in I . h is false in I . So this clause is false in I . Therefore I isn't a model of KB .

- Contradiction.

- The C generated at the end of the bottom-up algorithm is called a **fixed point**.
- Let I be the interpretation in which every element of the fixed point is true and every other atom is false.
- I is a model of KB .
Proof: suppose $h \leftarrow b_1 \wedge \dots \wedge b_m$ in KB is false in I . Then h is false and each b_i is true in I . Thus h can be added to C .
Contradiction to C being the fixed point.
- I is called a **Minimal Model**.

If $KB \models g$ then $KB \vdash g$.

- Suppose $KB \models g$. Then g is true in all models of KB .
- Thus g is true in the minimal model.
- Thus g is in the fixed point.
- Thus g is generated by the bottom up algorithm.
- Thus $KB \vdash g$.